

Finger Trees

The following exercises are executed in two steps. First (Exercises 1–5), we do a few sanity checks and paper exercises to help you understanding the article of Hinze and Paterson on finger trees. You may skip the first 5 exercises if you were comfortable with the notation in the paper.

After these introductory exercises, we proceed to implement Finger Trees in Scala to deepen your understanding of how the data structure works, and to prepare better for the exam.

Since there is no mandatory hand-in this week, we break with the course tradition of cramming all code into a single file (used purely to facilitate easy hand-in and grading). We keep types in separate files, and we separate Spec files as well, as more commonly done in Scala.

Our implementation is designed to be eager, following the regular strictness of Scala. However it would be an interesting exercise to extend it so that it is possibly lazy, like in the paper of Hinze and Paterson. The obvious choice is to make values of elements stored in the queue lazy. Then there is also a discussion of possible suspension of the middle element of the tree on page 7 (we do not explore this).

Hand-in: no hand-in this week.

Exercise 1. (somewhat trivial; skip if you know what a deque is)

Assume that we have a simple double-ended queue (a deque). We will make one diagram per each point below.

1. Draw how the dequeue looks after adding 1 on the left, and then adding 2 on the left, and then adding 3 on the left.
2. Now take the dequeue from the previous point and draw how it looks after adding 4 on the **right**, and then after adding 5 on the **right**.
3. Now take the dequeue resulting from the previous point and draw how it looks after popping two elements from the **left**, and one element from the **right**.

Exercise 2. Assume that we have an empty finger tree ready to store integer numbers. We perform the following operations on the tree. Draw how the tree looks after the series of operations in each point.

1. We add number 1, and then number 2 (both from the left).
2. Then add numbers 3, 4 and 5 from the left (in this order)
3. Then add number 6 from left
4. Then add numbers 7, 8 and 9 from the **left**.
5. Then add number 10 from the **right**.
6. Then remove a number from the **right** twice.
7. Then remove a number from the **right**.

Exercise 3. This exercise uses Haskell notation to stay close to the paper. So (+) denotes the binary plus operator and (:) denotes the cons operator for lists (in Scala denoted using double colon (_ : _) or Cons (_ , _)). Square brackets in Haskell are used to denote list literals, with empty square brackets representing the empty list (Nil or List() in Scala). The reduce right and reduce left functions are defined on page 3 in the paper.

What is the result of running:

1. `reducer (+) [1,2,3,4] 0`
2. `reducerl (+) 0 [1,2,3,4]`
3. `reducer (+) (Node3 1 2 3) 0`
4. `reducerl (+) 0 (Deep [1,2] (Single (Node3 3 4 5)) [0])`
5. `reducer (:) (Deep [1,2] (Single (Node3 3 4 5)) [0]) []`
6. `reducerl (:) [] (Deep [1,2] (Single (Node3 3 4 5)) [0])`

Exercise 4. Write the following Haskell data terms using the corresponding Scala type constructors from `FingerTree.scala`

1. `[1,2,3]`
2. `[]`
3. `Node2 1 2`

Observe that we work with `Node2[Int]` here (the type parameter is inferred).

4. `Node2 (Node2 1 2) (Node2 3 4)`

After solving, make type parameters to `Node2` explicit, to see whether you understand the types.

5. `Deep [1,2] (FingerTree (Node3 3 4 5)) [0]`

Also make the type parameters explicit.

Exercise 5. Translate to Scala the following Haskell type expressions. Use the algebraic data types for finger trees as defined in `FingerTree.scala`. Double colon in Haskell denotes a typing annotation, like a single colon in Scala. Haskell uses single arrows (`->`) as function type constructors, whereas Scala uses double arrows (`=>`). Small letters (free names) in Haskell types denote type variables. These can be both usual type parameters and higher kinded type parameters. Square brackets are also the list type constructor in Haskell, so `[a]` roughly corresponds to `List[A]` in Scala (adjusted for conventions).

NB. Some of these are solved in the Scala file, but please try to solve them yourself – then you will be able to work with the Scala file and with the paper more easily.

1. `addL :: a -> FingerTree a -> FingerTree a`
2. `addLprime :: f a -> FingerTree a -> FingerTree a`
3. `toTree :: f a -> FingerTree a`
4. `viewL :: FingerTree a -> ViewL FingerTree a`
5. `deepL :: [a] -> FingerTree a -> ViewL FingerTree a`

Exercise 6. Read Section 2.1 in Hinze/Paterson, and recall our own definition of the `Monoid` type class, and its instances. Just read through the provided file `Monoid.scala` (a copy from a few weeks back) to recall what this is all about.

Exercise 7. Study Section 2.2 in Hinze/Paterson. The file `Reduce.scala` implements the type class `Reduce` in Scala. Below the type class (implemented as a trait), there is a companion object that can be used to store instances of the type class for standard library types. Note that we implement foldable laws using the monoid instance from the previous exercise (to test whether the two reducers are the same in a monoid). This law is based on the argument in the opening of Section 2.2.

Implement the instance for `Reduce` in there. Check that the automatic tests are passing.

Exercise 8. The final part of Section 2.2 shows how to implement `toList` using an instance of `Reduce`. We could have implemented it the same way in Scala, placing a function `toList` in the `Reduce` companion object and requiring that it is run on type constructors for which an instance of `Reduce` exists. However, unlike Haskell, Scala is object oriented, so we can put `toList` inside the `Reduce` trait, when the existence of the instance is already guaranteed (`this`). This makes the access to the function a bit easier and more direct for users, so let do so. Complete this implementation in the marked place in the trait. Check that the automatic tests are passing.

Exercise 9. Read the opening of Section 3 (the first two pages, including the first page of Section 3.1). Open `Node.scala` and `FingerTree.scala` and find types `Node` and `FingerTree` implemented in Scala. Note that the type `Tree` is not implemented, as it is merely brought up as a motivating example in the section, and not used later in the paper.

Make sure to check not only the abstract trait, but also the case classes realizing the data types from Haskell. Finally, point to the location in Scala's type definitions where *polymorphic recursion* is visible in types.

No programming in this exercise.

Exercise 10. We are now moving to the final part of Section 3.1 (page 5) with the goal of implementing an instance of the `reduce` type class for nodes. We place it in the companion object of `Node` towards the end of the file (`Node.scala`). Implement this type class instance providing definitions for `reduceR` and `reduceL` for nodes. Make sure that the automated tests pass.

Exercise 11. Implement reducers for finger trees as specified in Haskell in the final part of Section 3.1 (page 5). We place the type class instance (of `Reduce`) for finger trees in the companion object `FingerTree`, in the place marked with the exercise number in `FingerTree.scala`. Note that you will have to use the instances of `Reduce` type class for nodes and digits when solving this exercise.

Aim for the solution to compile. We develop some tests much later, in exercises 16–17.

Exercise 12. Read the first half page of Section 3.2, which specifies how to add an element to the left of the deque, implemented in a finger tree. Note that in the paper the name of the function is \triangleright , and that it is written infix, so the first argument is to the left of the triangle. Name the function `addL`.

After the tests for `addL` pass, implement also `addR`. We develop tests for `addR` in Exercise 13.

The implementations of `addL` and `addR` should involve *polymorphic recursion*. Find (and mark with a comment) the calls which are polymorphically recursive.

Exercise 13. Write three tests for `addR`:

- A scenario test that adds 42 to an empty tree, then converts the tree to a list and checks whether this list is a singleton containing 42.
- A scenario test that adds 42 to an empty tree, then adds 42, and converts the tree to a list and checks whether this list contains 42 and 43.
- Generalize the above tests to a property test: for any list `l` adding elements from the left (from the head) with `addR` will produce a tree that contains the same list (for instance after converting it back to a list)

Of course, your tests should all pass. If they fail, search for bugs either in the test or in your implementation of `addR`.

Exercise 14 [optional]. We named the functions `addL` and `addR`. If you like some syntactic fun, use a unicode triangle as an alias. Add a method named `<` to trait `FingerTree` and use it in later exercises instead of `addR`. Note that `<` can be called infix, just like in the Haskell examples.

Adding `addL` is a bit more difficult, as the left argument needs to be any type `A`—we cannot simply add it to `FingerTree`. Add an extension method (named `>`) for any type `A`, using implicit conversions (pimp-up-my-library pattern), to be able to call `addL` infix.

Exercise 15. We skip the lifted (primed) versions of `addR` and `addL` and proceed to implement `toTree` from page 6. You can use `reduce` directly in `toTree` to save time on implementing the primed versions.

Exercise 16. Implement the instance of the type class `Gen` for `FingerTree`, so that we can use trees in property tests. Do it in two steps. First, implement `fingerTreeOfN[A] (n: Int)` that creates a tree containing `n` elements. You can use `toTree` in this implementation. Second, create `fingerTree[A]` by invoking `fingerTreeOfN` with a generated size. Be careful to limit the sizes, to not run out memory.

Finally, expose the latter generator as an implicit instance of `Arbitrary` for `FingerTree`.

We will use these generators in tests for `Reduce[FingerTree]` in the next exercise.

Exercise 17. Implement the property tests for `Reduce[FingerTree]` using `Reduce.Law`. You can get inspiration from tests for `Reduce[List]` and `Reduce[Node]`. Can you find them?

Exercise 18. Find the definition of type `ViewL` in the paper, and then the corresponding Scala implementation `ViewL` in `FingerTree.scala`. Convince yourself that you understand the proposal for Scala. No coding in this exercise.

The paper uses explicit views to implement deconstruction of finger tree values (removing elements with pattern matching). Read the provided reading material from Horstmann about extractors in Scala and observe that they could be used for the same purpose like views, but more elegantly than in the Haskell examples. With Scala's extractors, we do not have to call the view function explicitly. Thus instead of creating an ADT for views, we will define three extractor objects (in the next exercise).

To check your understanding, find the implementation of the extractor `Digit` in the object `Digit`. Find the pattern matching on `Digit` elsewhere in the code that used this extractor.

Exercise 19 [difficult]. Implement extractors `Nil`, `ConsL`, and `ConsR`. The behavior of `Nil` is specified in the first case of definition of `viewL` in the paper. The behavior of `ConsL` is specified by the two remaining cases. Behaviour of `ConsR` will be symmetric.

Note that this exercise requires also implementing `deepL` and `deepR`, presented on the next page in the paper. The tests use the `head` and `tail` functions defined in page 7, and already implemented using your pattern extractors in the trait `FingerTree`. You can check how patterns are used in their code in the trait.

Our finger tree types also implement `empty` and `nonEmpty`, but (page 7) but without using pattern matching and views. We exploit dynamic dispatch and inheritance in Scala instead (check!). This is much simpler, in our humble opinion, but Haskell has no dynamic dispatch, so the authors did not have that option.